[AI System Design – Term Project]
Optimizations of AI Accelerator Networks

Chester Sungchung Park (박성정)
SoC Design Lab, Konkuk University
Webpage: http://soclab.konkuk.ac.kr

Teaching Assistants

- Shinyoung Kim (shinyoungkim@konkuk.ac.kr)
- Jaesuk Lee (jaesuki98@konkuk.ac.kr)

Outline

- Introduction
  - Target workload
  - NetTLMSim
- Problem to solve
- Reference codes
- Evaluation
- Submission

Target Workload

- BERT-medium ***Layer group A***
  - Layer group 0 + matmul (QK^T)

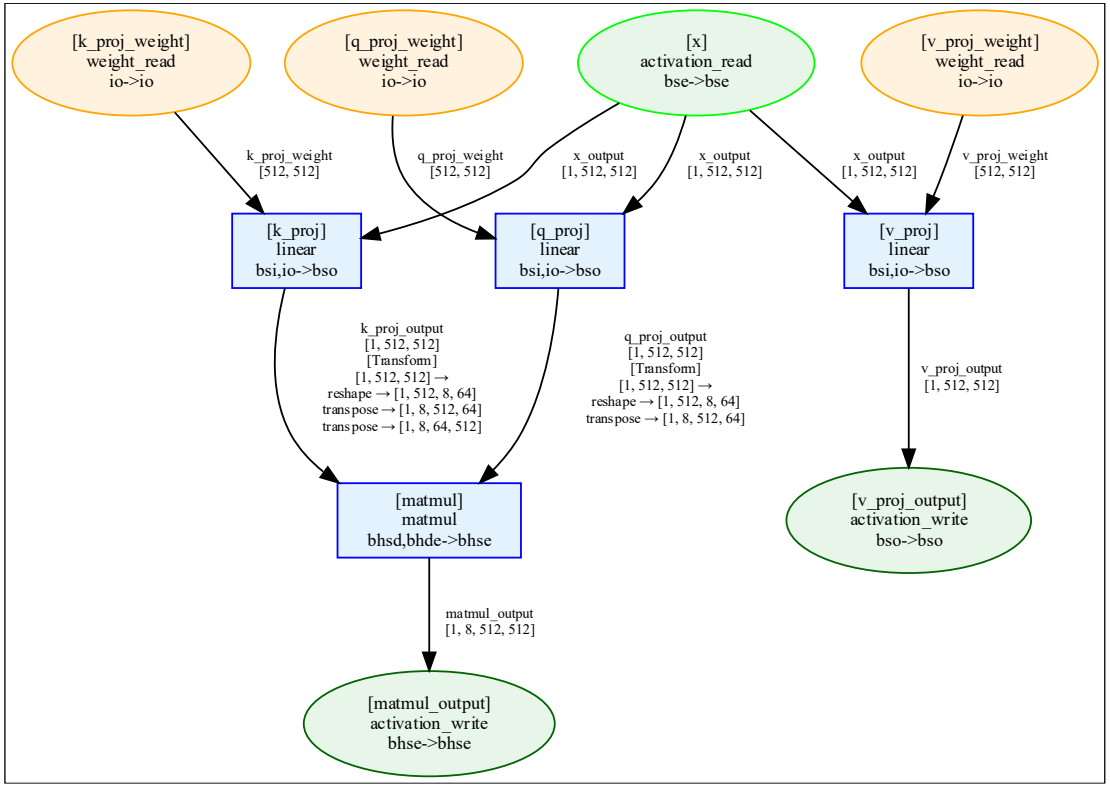

NetTLMSim

❑ Virtual prototype

Assumed design parameters

Space-Time Mapping	Target Hardware	16 chiplets per package [1], 16 PEs per chiplet [1, 2], 1024 MACs per PE [1, 3] Total 512 TOPS at 1 GHz freq.
	PE Style	NVDLA-like vector MAC array
	Loop Order	OS-LWS [4]
Network Architecture	Topology	Mesh [2]
	Routing	DOR YX [2]
	Flow Ctrl	Cut-through [2]
	Packet Len	1 head flit, 16 body flits [2]

Memory Architecture	DRAM Type	HBM2E, 2 devices Total 1024 GB/s
	SRAM Type	Dual port [3]
	Global Buf. Size, Bandwidth	640 KB activation storage [4], 128 GB/s per NoC router
	Local Buf. Size, Bandwidth	36 KB x2 (double buffered) [4], 128 GB/s at 1 GHz freq.
	Precision	8 bits (24 bits for partial sums) [4]

[1] Cai at el. 2024
[2] Zimmer at el. 2020
[3] Keller at el. 2023
[4] Venkatesan at el. 2019

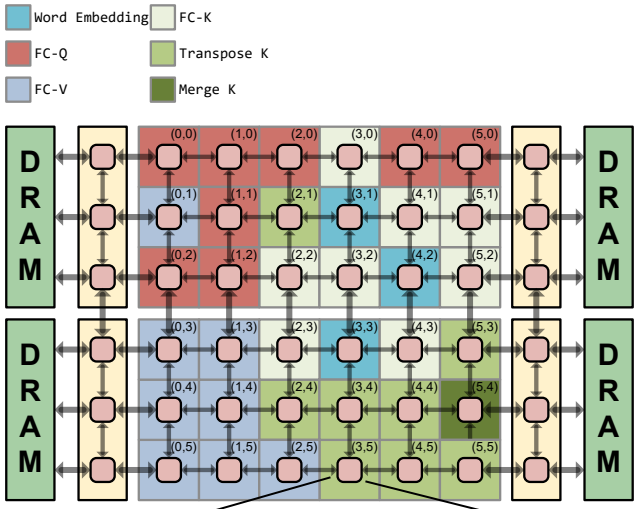

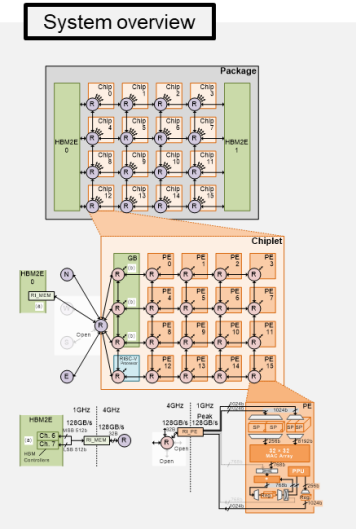

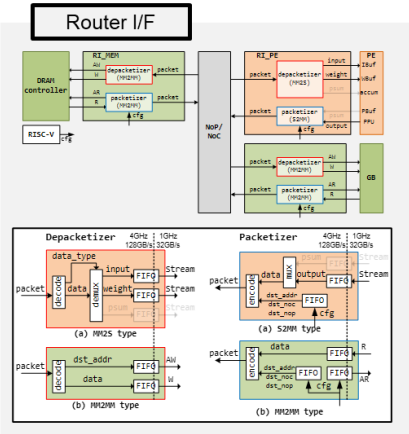

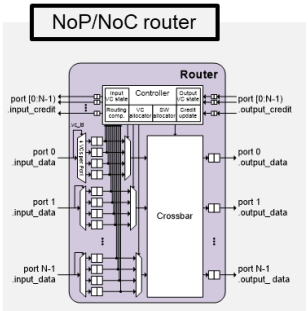

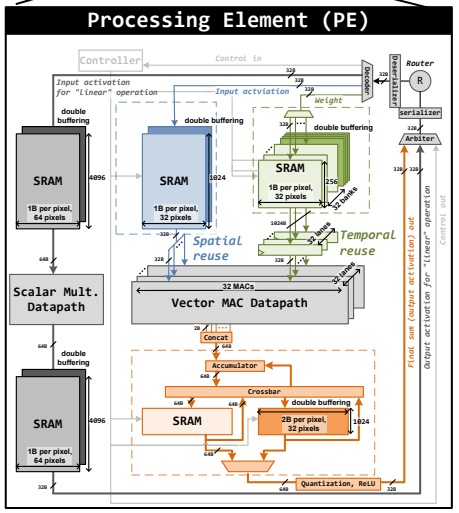

For more details, check the following presentation:
F. S. Park and C. S. Park, *Pre-RTL simulation based design space exploration for multi-chiplet dataflows*, HiPChips (MICRO 2024).

NetTLMSim

- Virtual prototype (cont'd)

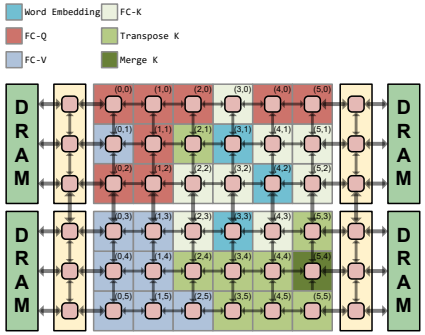

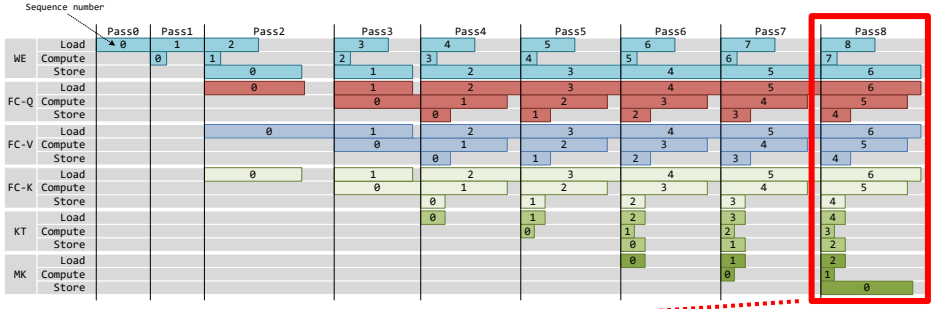

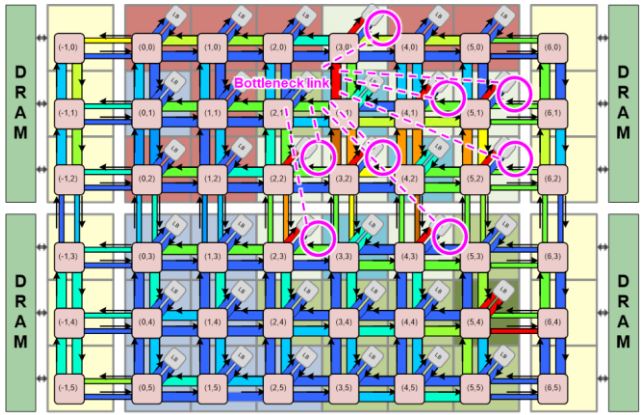

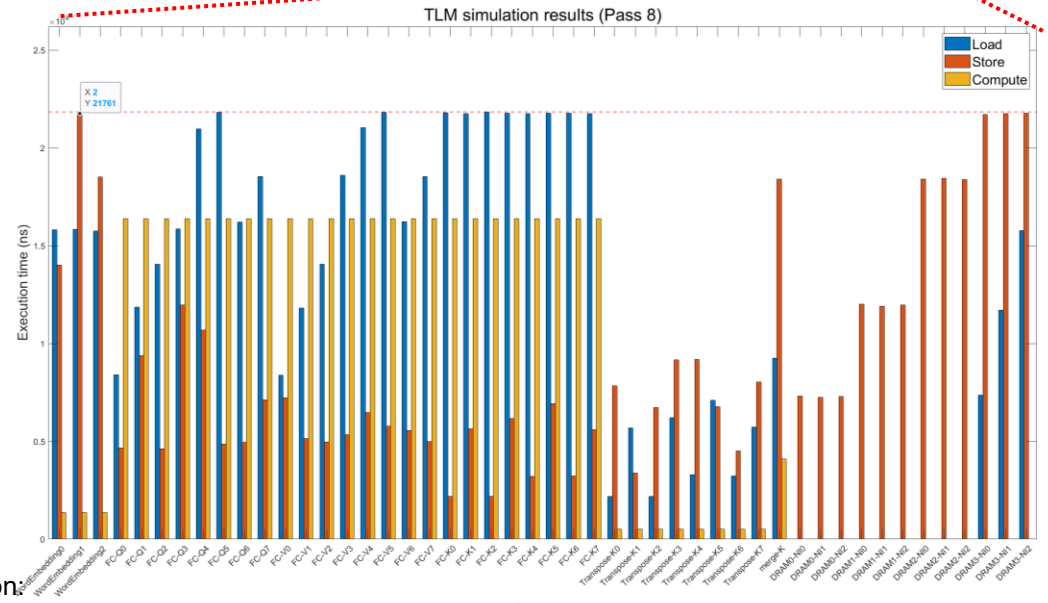

For more details, check the following presentation:
F. S. Park and C. S. Park, *Pre-RTL simulation based design space exploration for multi-chiplet dataflows*, HiPChips (MICRO 2024).

Reference Project

- Build project and run simulation (See ***Appendix A***)

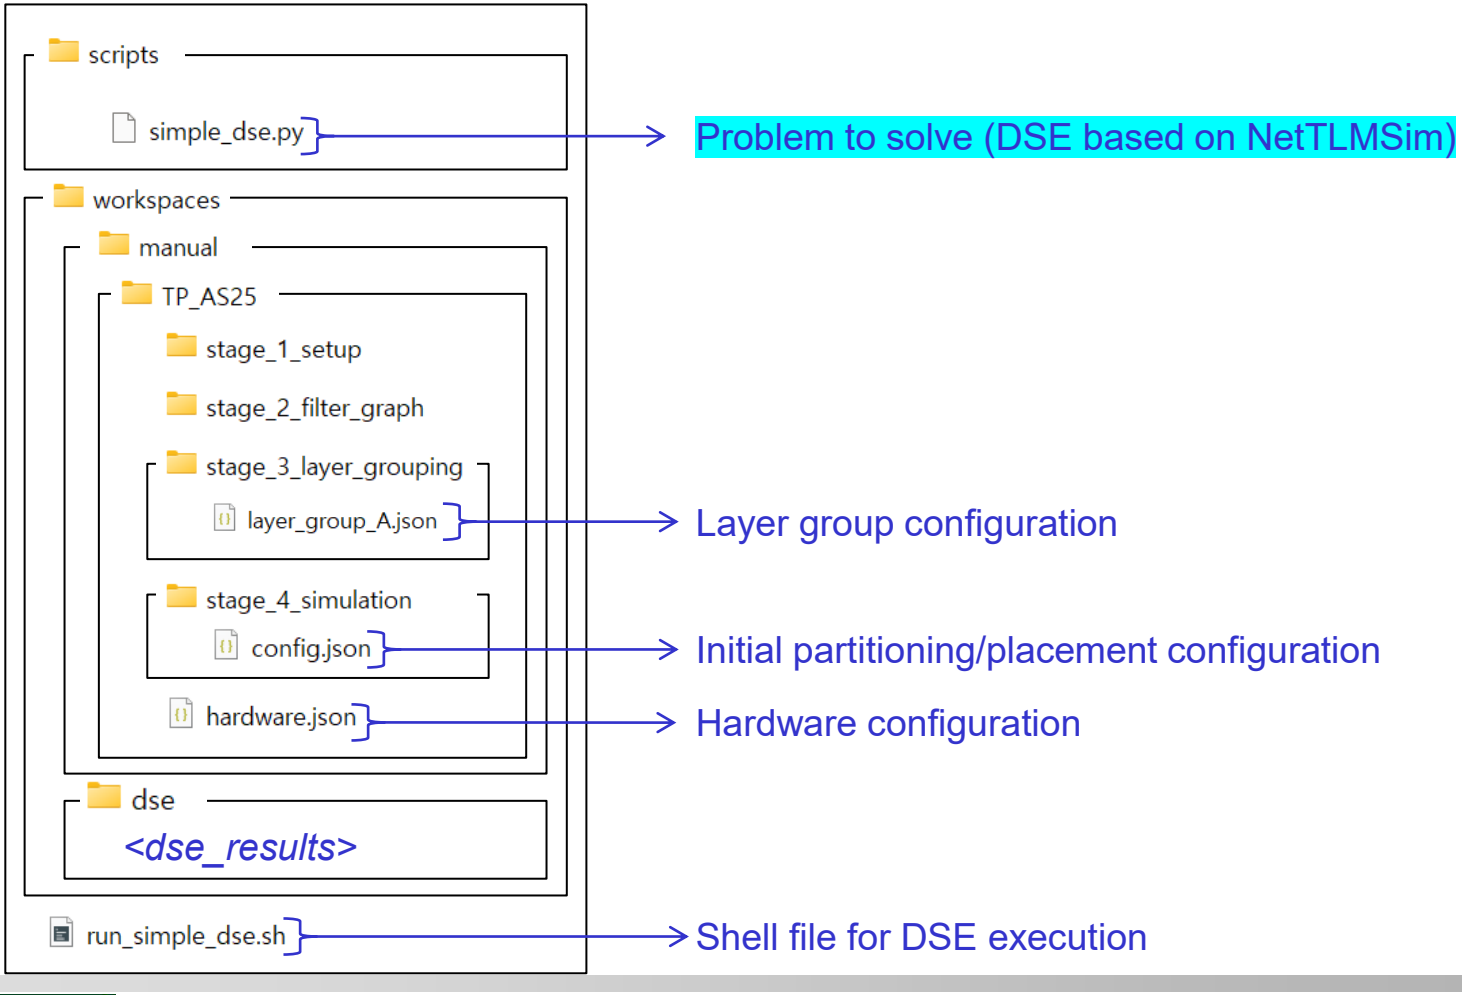

Problem to Solve

- Partitioning/placement configuration selection
  - Input: **NetTLMSim output** for the *n*-th iteration
  - Output: **NetTLMSim input** for the (*n*+1)-th iteration
    - ✓ Only partitioning/placement configuration *varied*
    - ✓ Hardware/layer group configuration *kept fixed*

Problem to Solve

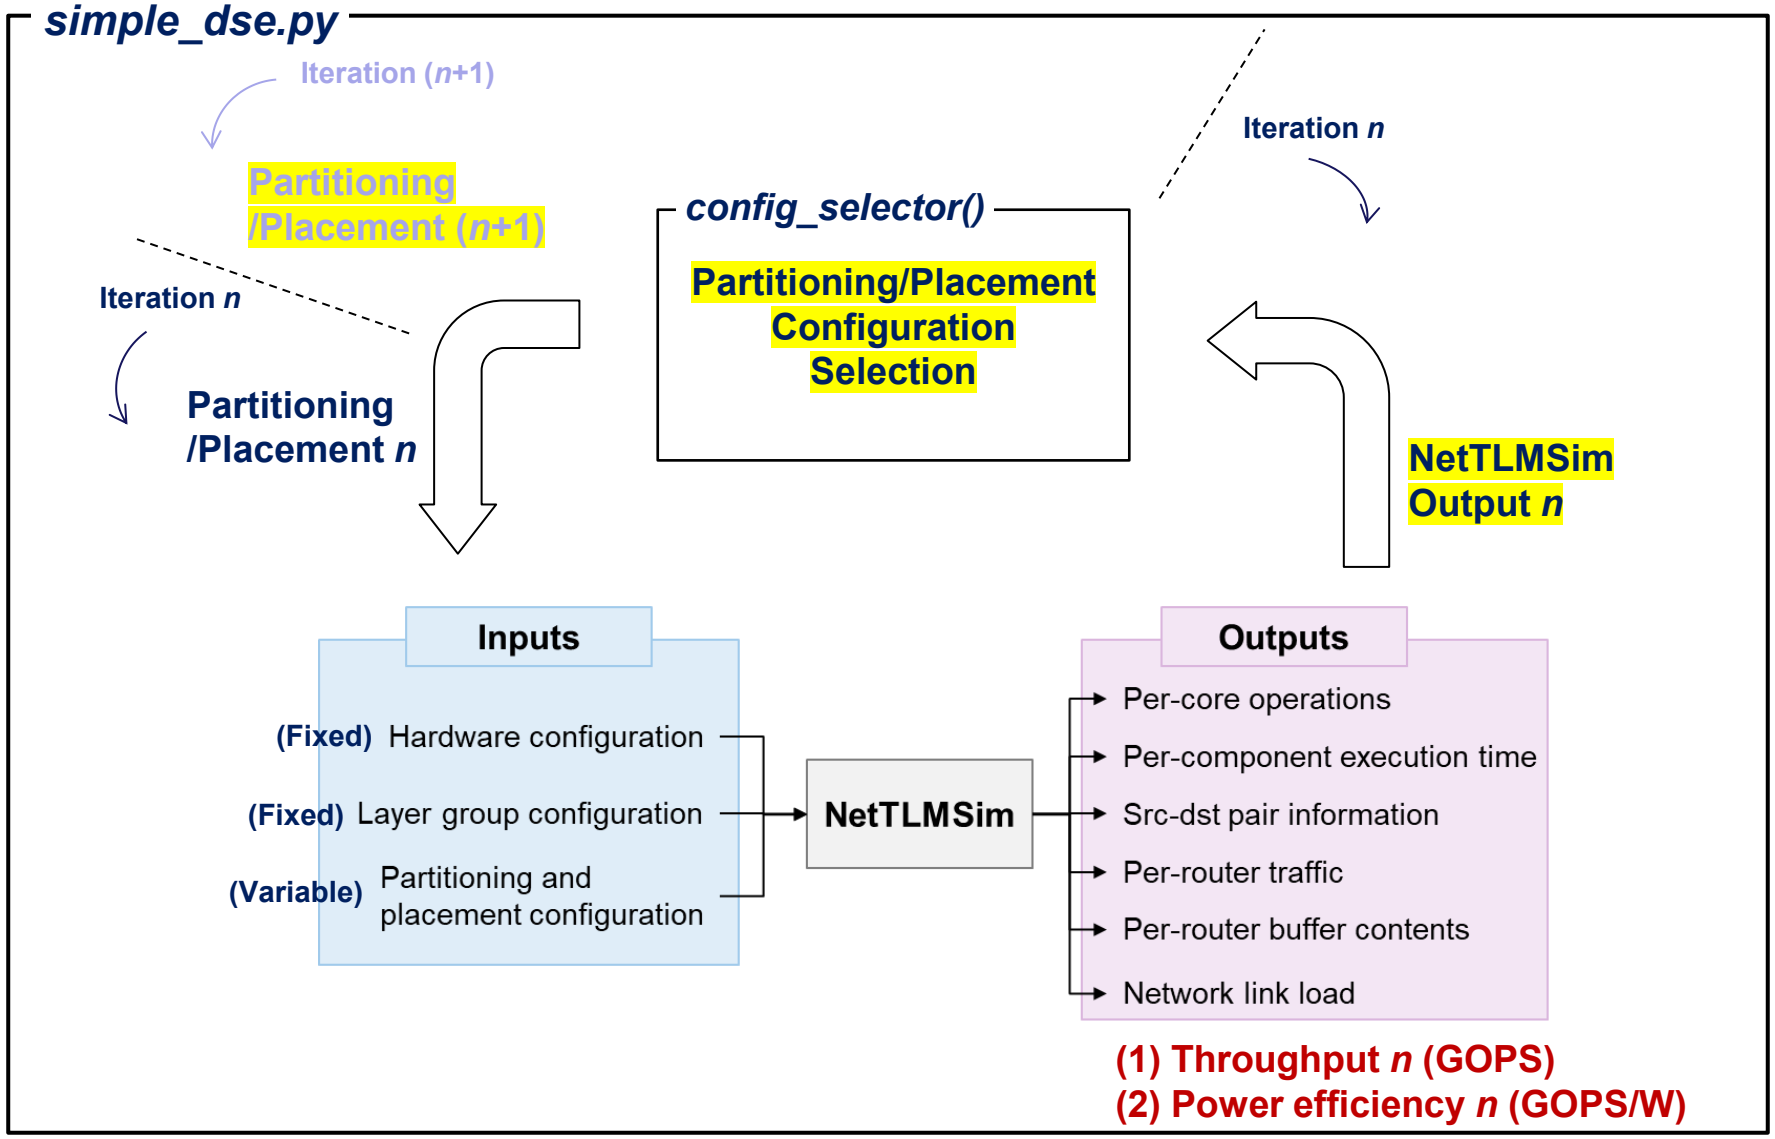

Problem to Solve

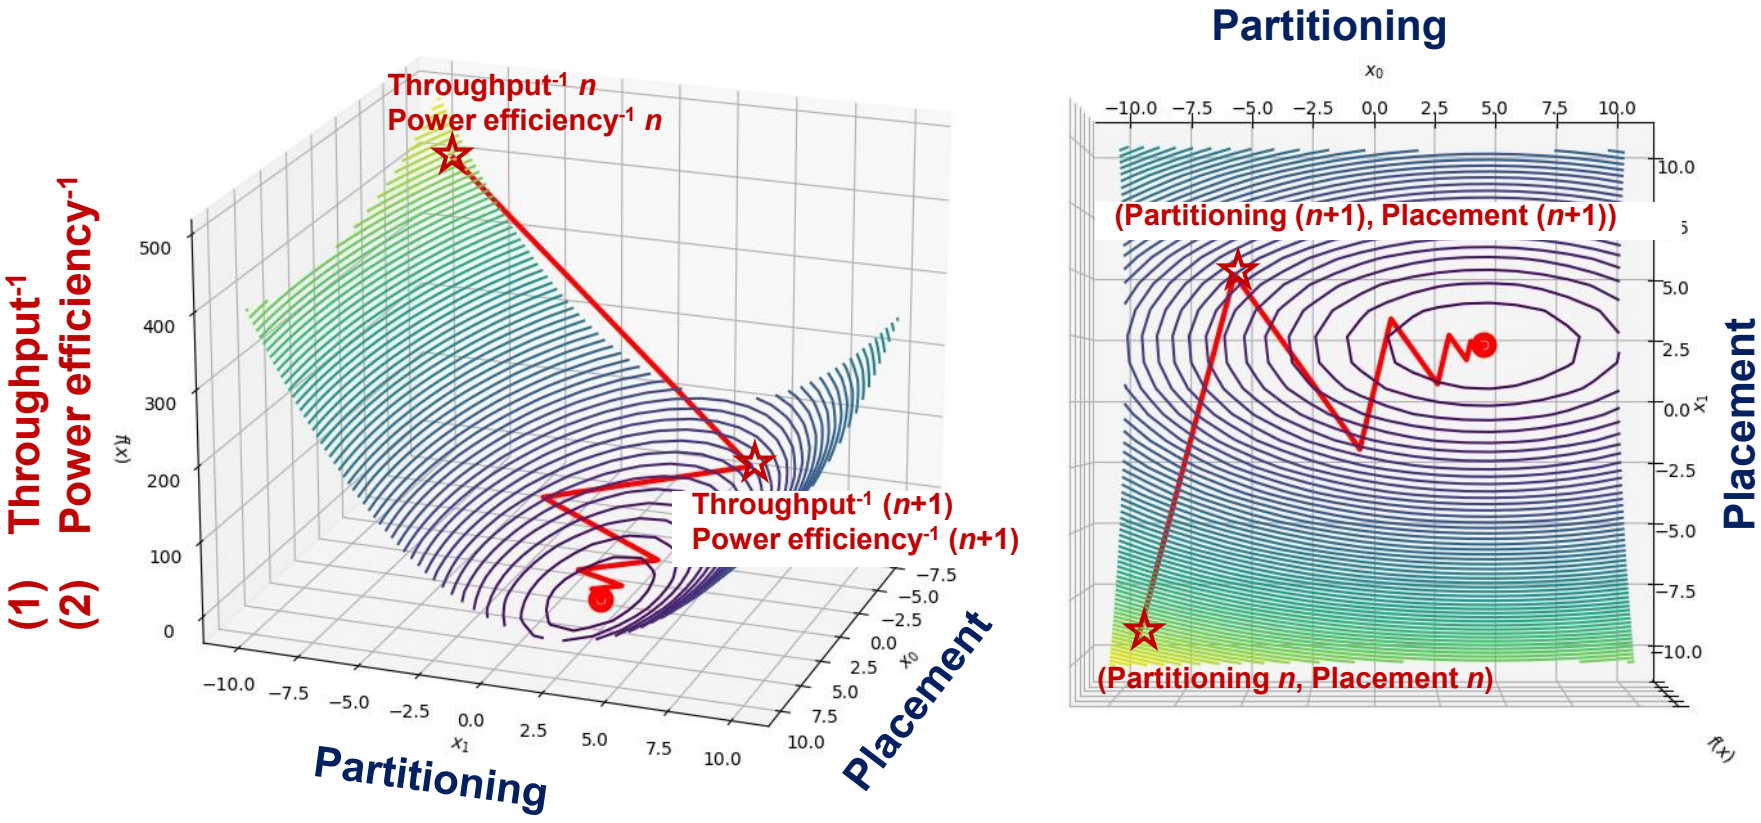

Design Constraints

- Optimize with respect to partition and placement only
  - You are ***not*** allowed to modify any other settings, e.g., hardware, layer group configurations etc.
    - ✓ Otherwise, it won't be considered for evaluation
- Run *no* additional simulations (e.g., those using NetTLMSim) *inside* the partitioning/placement configuration selection
  - Only a single simulation run is allowed per iteration.

Design Constraints

- Modify **only** (a part of) the body of ***simple_dse.py***
  - You are allowed to change ***only*** the paragraphs between ***"Edit code below"*** and ***"Edit code above"***

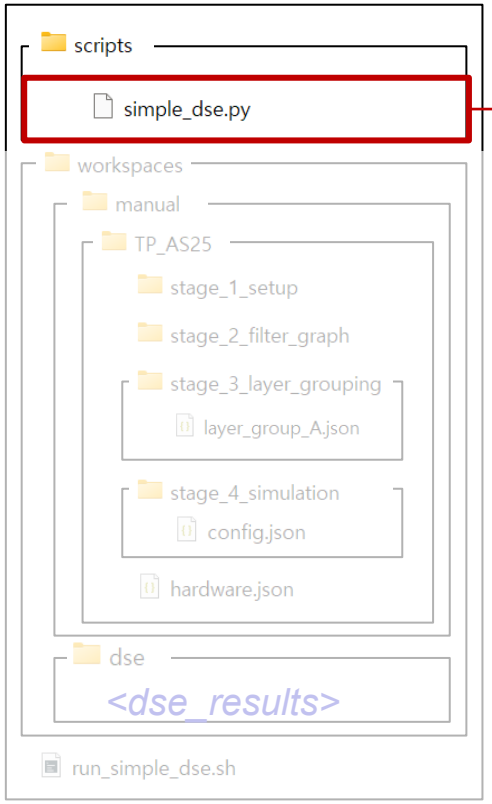


```python
def config_selector(
    config: Dict[str, Any],
    iteration: int,
    allowed_groups: Optional[List[str]] = None,
    last_iter_dir: Optional[Path] = None,
) -> Dict[str, Any]:
    partitions = config.get("partitions")
    if not isinstance(partitions, dict):
        return config
    # Load last iteration logs
    if last_iter_dir.exists() & (iteration != 1):
        ...

    """ Edit code below """
    def _rotate_router_ids(
        nodes: Dict[str, Any],
        skip_if: Optional[Callable[[str, Dict[str, Any]], bool]] = None,
    ) -> None:
        """Mutate nodes in-place by rotating their `router_ids`."""
        for node_key, node in nodes.items():
            if skip_if is not None and skip_if(node_key, node):
                continue

            ids = node.get("router_ids")
            if isinstance(ids, list) and ids:
                if len(ids) == 1:
                    shift = 1
                else:
                    shift = random.randint(1, len(ids) - 1)
                node["router_ids"] = ids[shift:] + ids[:shift]

    for group_key, group in partitions.items():
        try:
            group_idx = group_key.split("_")[-1]
        except (ValueError, IndexError):
            group_idx = None

        if allowed_groups is not None and group_idx not in allowed_groups:
            continue

        compute_nodes = group.get("compute_nodes") or {}
        _rotate_router_ids(compute_nodes)

        tensor_nodes = group.get("tensor_nodes") or {}
        _rotate_router_ids(
            tensor_nodes,
            skip_if=lambda node_key, node: "weight" in node_key,
        )

    """ Edit code above """
    # Ensure mutated partitions stay attached
    config["partitions"] = partitions

    # Record a breadcrumb so users can see what changed
    meta = config.setdefault("_simple_dse", {})
    history = meta.get("history", [])
    history.append(
        {
            "iteration": iteration,
            "timestamp": datetime.now().isoformat(),
            "target_groups": allowed_groups if allowed_groups is not None else "all",
        }
    )
    meta["last_iteration"] = iteration

    return config
```

```python
if last_iter_dir.exists() & (iteration != 1):
    log_dir_path = last_iter_dir / "logs/layer_group_A"
    sim_result_path = last_iter_dir / "simulation_result.json"
    log_alloc_path = log_dir_path / "allocation.csv"
    log_exec_time_path = log_dir_path / "execution_time.csv"
    log_flowID_path = log_dir_path / "flowID.csv"
    log_link_load_path = log_dir_path / "link_load.csv"
    log_routers_path = list(log_dir_path.glob("router_*.csv"))
    log_router_path = [f for f in log_routers_path if re.match(r'^router_\d+\.csv$', f.name)]
    log_router_buf_path = [f for f in log_routers_path if re.match(r'^router_\d+_buf\.csv$', f.name)]
    log_d2d_buf_path = list(log_dir_path.glob("d2d_*_*_buf.csv"))

    with open(sim_result_path, "rb") as f:
        sim_result = orjson.loads(f.read())
    throughput = sim_result['data']['summary']['total_throughput_ops']
    power_effs = sim_result['data']['summary']['power_efficiency_ops_per_w']
    alloc = pd.read_csv(log_alloc_path)
    exec_time = pd.read_csv(log_exec_time_path)
    flowID = pd.read_csv(log_flowID_path)
    link_load = pd.read_csv(log_link_load_path)
    routers = [pd.read_csv(file, skipinitialspace=True) for file in log_router_path]
    router_bufs = [pd.read_csv(file, dtype=str, skipinitialspace=True) for file in log_router_buf_path]
    d2d_bufs = [pd.read_csv(file, dtype=str, skipinitialspace=True) for file in log_d2d_buf_path]
```

Using these variables to utilize the results from the previous iteration

Evaluation

- Submission completeness (10pt)
  - Reproducibility
- Optimality (40pt)
  - Final throughput and power efficiency
- DSE efficacy (30pt)
  - Sampling efficiency
  - General applicability
- Documentation (20pt)
  - Covering all those mentioned in the above

Submission Completeness

- Make sure that your submission is **complete**
  - In other words, it should be possible to reproduce your design together with the claimed throughput and power efficiency using **only** the files that you submitted by the submission deadline

Optimality

- Evaluated both absolutely and relatively
  - In other word, the quantity as well as the ranking matters
- Evaluate separately for throughput and power efficiency

DSE Efficacy

- Sampling efficiency
  - Defined as the final quantity divided by the total number of iterations
    - ✓ Evaluates *how many simulation runs are required to reach the final throughput and power efficiency*
  - Considered for **both** *training* (fine-turning) and *inference* in the case of AI/ML (and separately, if necessary)
  - The fewer iterations, the better sampling efficiency
- General applicability
  - Evaluates whether the proposed idea is generally applicable
    - ✓ For example, different layer group configurations (including layer groups 0~2)
  - May be further verified by extra hardware/layer group configurations

Submission

- Deadline
  - **Dec. 12 (Fri), 10:00** GMT+9
- **Only one zip file** submission **per team** including the following files:
  - Source code (all those needed to reproduce your results)
    - ✓ Except all those provided by TA (e.g., reference project)
  - Documentation (PPT)
    - ✓ Including the explanation of the above source code
- Upload the zip file to the Ecampus
  - Plus, send to chesterku2013@gmail.com as a *backup*
- You can post questions in the Ecampus (Q&A)
- Delayed submission will result in penalty!

Appendix A:
Running a DSE

Runing a DSE

❑ Start the DSE

- Open the terminal with ***'Ctrl + `'***
- Type the following command and press ***'Enter'***
  - ✓ `./run_simple_dse.sh TP_AS25 --layer-groups A --iterations` ***N***
  - ✓ Set iteration count to ***N***